

Exercise 8: Clustering stocks using KMeans

In this exercise, you'll cluster companies using their daily stock price movements (i.e. the dollar difference between the closing and opening prices for each trading day). You are given a NumPy array `movements` of daily price movements from 2010 to 2015, where each row corresponds to a company, and each column corresponds to a trading day.

Some stocks are more expensive than others. To account for this, include a `Normalizer` at the beginning of your pipeline. The `Normalizer` will separately transform each company's stock price to a relative scale before the clustering begins.

Normalizer vs StandardScaler

Note that `Normalizer()` is different to `StandardScaler()`, which you used in the previous exercise. While `StandardScaler()` standardizes **features** (such as the features of the fish data from the previous exercise) by removing the mean and scaling to unit variance, `Normalizer()` rescales **each sample** - here, each company's stock price - independently of the other.

This dataset was obtained from the Yahoo! Finance API.

From the course *Transition to Data Science*. [Buy the entire course for just \\$10](#) for many more exercises and helpful video lectures.

Step 1: Load the data (*written for you*)

```
In [2]: import pandas as pd
```

```
fn = '../datasets/company-stock-movements-2010-2015-incl.csv'
stocks_df = pd.read_csv(fn, index_col=0)
```

Step 2: Inspect the first few rows of the DataFrame `stocks_df` by calling its `head()` function.

```
In [6]: stocks_df.head()
```

```
Out [6]:
```

	2010-01-04	2010-01-05	2010-01-06	2010-01-07	2010-01-08	\
Apple	0.580000	-0.220005	-3.409998	-1.170000	1.680011	
AIG	-0.640002	-0.650000	-0.210001	-0.420000	0.710001	
Amazon	-2.350006	1.260009	-2.350006	-2.009995	2.960006	
American express	0.109997	0.000000	0.260002	0.720002	0.190003	
Boeing	0.459999	1.770000	1.549999	2.690003	0.059997	
	2010-01-11	2010-01-12	2010-01-13	2010-01-14	2010-01-15	\
Apple	-2.689994	-1.469994	2.779997	-0.680003	-4.999995	
AIG	-0.200001	-1.130001	0.069999	-0.119999	-0.500000	
Amazon	-2.309997	-1.640007	1.209999	-1.790001	-2.039994	
American express	-0.270001	0.750000	0.300004	0.639999	-0.130001	

Boeing	-1.080002	0.360000	0.549999	0.530002	-0.709999	
...		2013-10-16	2013-10-17	2013-10-18	2013-10-21	\
Apple	...	0.320008	4.519997	2.899987	9.590019	
AIG	...	0.919998	0.709999	0.119999	-0.480000	
Amazon	...	2.109985	3.699982	9.570008	-3.450013	
American express	...	0.680001	2.290001	0.409996	-0.069999	
Boeing	...	1.559997	2.480003	0.019997	-1.220001	
		2013-10-22	2013-10-23	2013-10-24	2013-10-25	2013-10-28
Apple	-6.540016	5.959976	6.910011	-5.359962	0.840019	
AIG	0.010002	-0.279998	-0.190003	-0.040001	-0.400002	
Amazon	4.820008	-4.079986	2.579986	4.790009	-1.760009	
American express	0.100006	0.069999	0.130005	1.849999	0.040001	
Boeing	0.480003	3.020004	-0.029999	1.940002	1.130005	
		2013-10-29				
Apple	-19.589981					
AIG	0.660000					
Amazon	3.740021					
American express	0.540001					
Boeing	0.309998					

[5 rows x 963 columns]

Step 3: Extract the NumPy array `movements` from the DataFrame and the list of company names (*written for you*)

```
In [14]: companies = list(stocks_df.index)
         movements = stocks_df.values
```

Step 4: Make the necessary imports:

- `Normalizer` from `sklearn.preprocessing`.
- `KMeans` from `sklearn.cluster`.
- `make_pipeline` from `sklearn.pipeline`.

```
In [7]: from sklearn.preprocessing import Normalizer
         from sklearn.cluster import KMeans
         from sklearn.pipeline import make_pipeline
```

Step 3: Create an instance of `Normalizer` called `normalizer`.

```
In [8]: normalizer = Normalizer()
```

Step 4: Create an instance of `KMeans` called `kmeans` with 14 clusters.

```
In [15]: kmeans = KMeans(n_clusters=14)
```

Step 5: Using `make_pipeline()`, create a pipeline called `pipeline` that chains `normalizer` and `kmeans`.

```
In [16]: pipeline = make_pipeline(normalizer, kmeans)
```

Step 6: Fit the pipeline to the `movements` array.

```
In [17]: pipeline.fit(movements)
```

```
Out[17]: Pipeline(steps=[('normalizer', Normalizer(copy=True, norm='l2')), ('kmeans', KMeans(n_clusters=14, n_init=10, n_jobs=1, precompute_distances='auto', random_state=None, tol=0.0001, verbose=0))])
```

In the next exercise: Let's check out your clustering!